

Machine Learning

Introduction

Welcome

Apple - iPhoto - New full-Sc X

www.apple.com/ilife/iphoto/

Store Mac iPod iPhone iPad iTunes Support

iLife '11

iPhoto iMovie GarageBand Video Showcase Resources Upgrade Now

 iPhoto '11

From your Facebook Wall to your coffee table to your best friend's inbox (or mailbox). Do more with your photos than you ever thought possible. And do it all in one place. iPhoto.

 Watch the iPhoto video ▶



What's New in iPhoto What is iPhoto?

Andrew Ng



Machine Learning

- Grew out of work in AI
- New capability for computers

Examples:

- Database mining

Large datasets from growth of automation/web.

E.g., Web click data, medical records, biology, engineering

- Applications can't program by hand.

E.g., Autonomous helicopter, handwriting recognition, most of Natural Language Processing (NLP), Computer Vision.

Machine Learning

- Grew out of work in AI

- Examples

- Examples



ig
host of

Machine Learning

- Grew out of work in AI
- New capability for computers

Examples:

- Database mining

Large datasets from growth of automation/web.

E.g., Web click data, medical records, biology, engineering

- Applications can't program by hand.

E.g., Autonomous helicopter, handwriting recognition, most of Natural Language Processing (NLP), Computer Vision.

Machine Learning

- Grew out of work in AI
- New capability for computers

Examples:

- Database mining

Large datasets from growth of automation/web.

E.g., Web click data, medical records, biology, engineering

- Applications can't program by hand.

E.g., Autonomous helicopter, handwriting recognition, most of Natural Language Processing (NLP), Computer Vision.

- Self-customizing programs

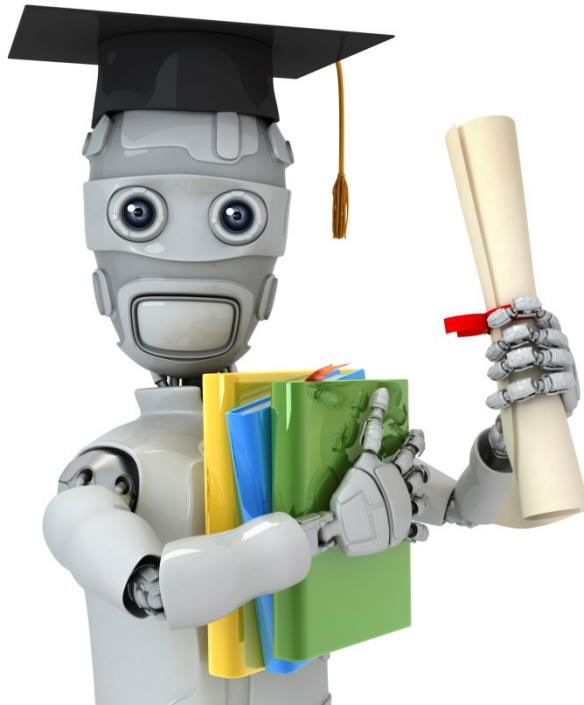
E.g., Amazon, Netflix product recommendations

Machine Learning

- Grew out of work in AI
- New capability for computers

Examples:

- Database mining
 - Large datasets from growth of automation/web.
E.g., Web click data, medical records, biology, engineering
- Applications can't program by hand.
 - E.g., Autonomous helicopter, handwriting recognition, most of Natural Language Processing (NLP), Computer Vision.
- Self-customizing programs
 - E.g., Amazon, Netflix product recommendations
- Understanding human learning (brain, real AI).



Machine Learning

Introduction

What is machine learning

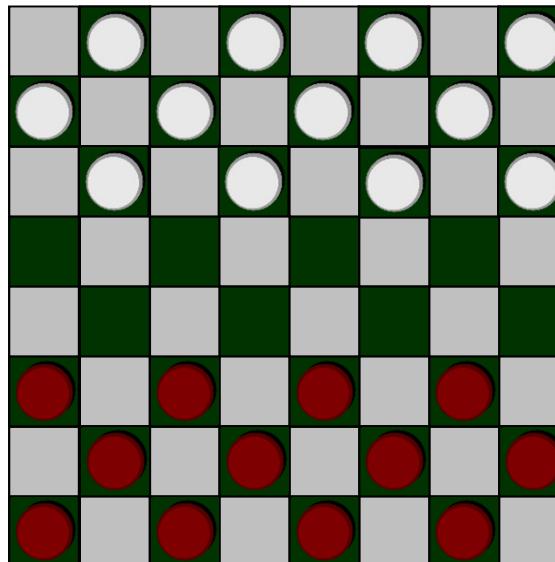
Machine Learning definition

Machine Learning definition

- Arthur Samuel (1959). Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.

Machine Learning definition

- Arthur Samuel (1959). Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.



Machine Learning definition

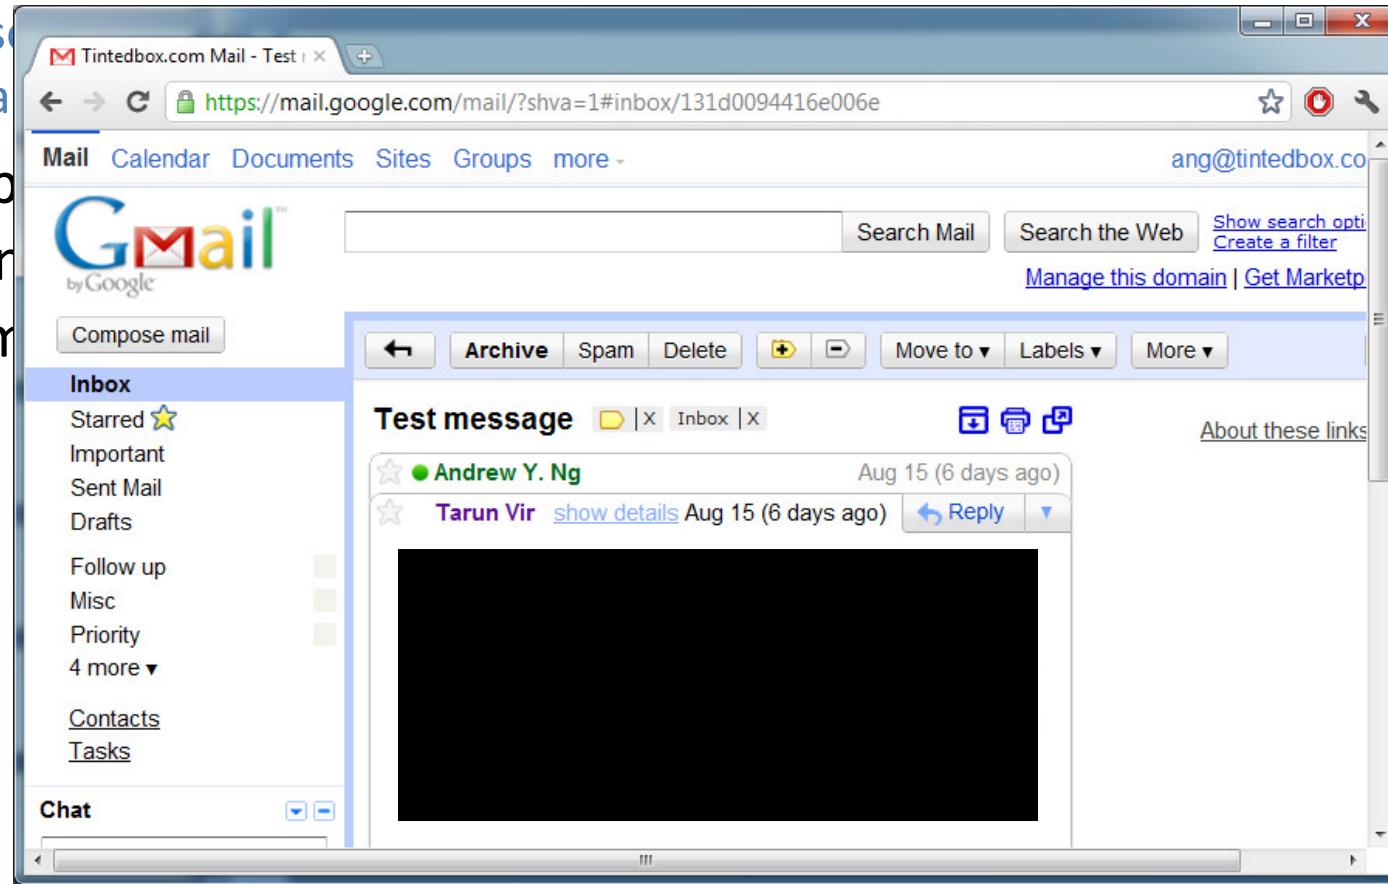
- Arthur Samuel (1959). Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.
- Tom Mitchell (1998) Well-posed Learning Problem: A computer program is said to *learn* from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E.

“A computer program is said to *learn from* experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E.”

Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam. What is the task T in this setting?

- Classifying emails as spam or not spam. *T ↙*
- Watching you label emails as spam or not spam. *E ↙*
- The number (or fraction) of emails correctly classified as spam/not spam. *P ↙*
- None of the above—this is not a machine learning problem. *P ↙*

“A computer program is said to *learn* from experience E with respect to



“A computer program is said to *learn from* experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E.”

Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam. What is the task T in this setting?

- Classifying emails as spam or not spam. *T ↙*
- Watching you label emails as spam or not spam. *E ↙*
- The number (or fraction) of emails correctly classified as spam/not spam. *P ↙*
- None of the above—this is not a machine learning problem. *P ↙*

Machine learning algorithms:


- Supervised learning
- Unsupervised learning



Others: Reinforcement learning, recommender systems.

Also talk about: Practical advice for applying learning algorithms.

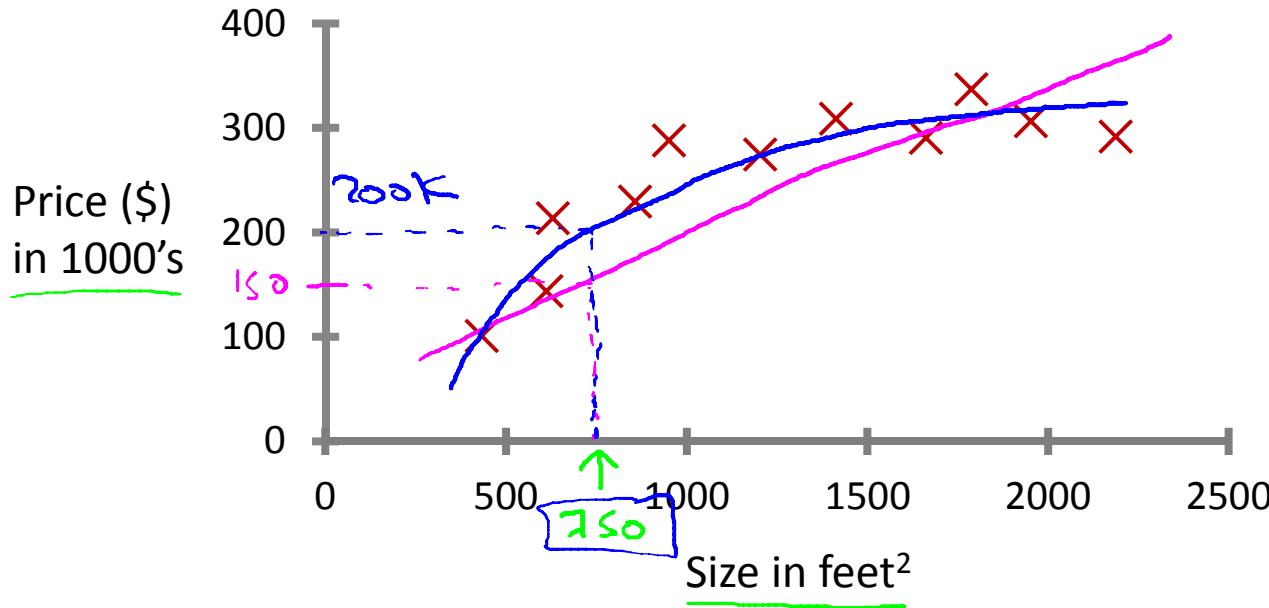




Machine Learning

Introduction Supervised Learning

Housing price prediction.

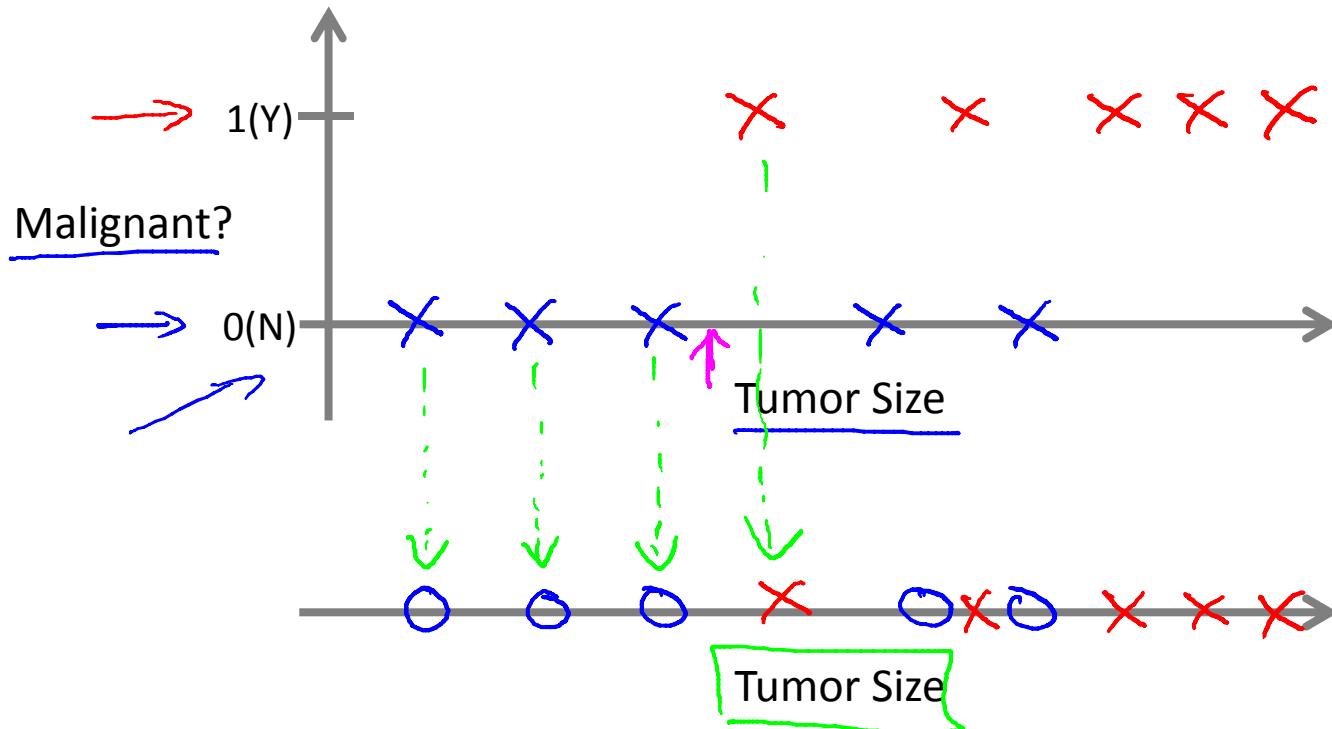


Supervised Learning

'right answers' given

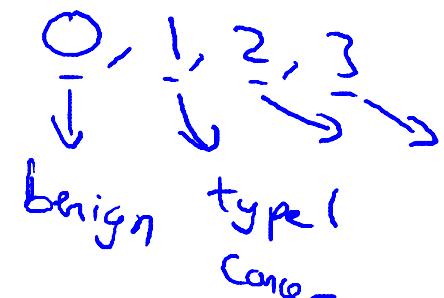
Regression: Predict continuous valued output (price)

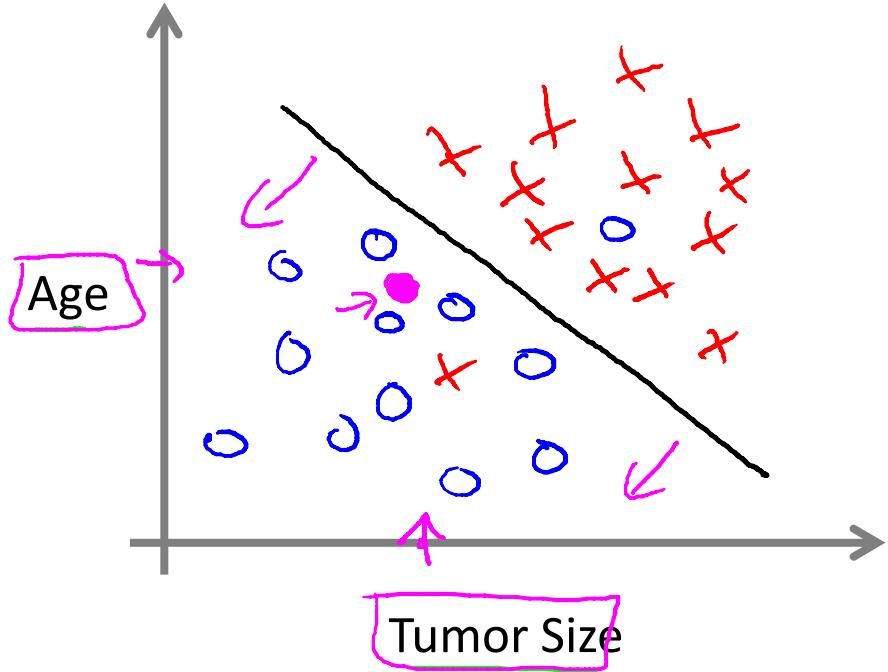
Breast cancer (malignant, benign)



Classification

Discrete valued output (0 or 1)





- Clump Thickness
- Uniformity of Cell Size
- Uniformity of Cell Shape
- ...

You're running a company, and you want to develop learning algorithms to address each of two problems.

1000's

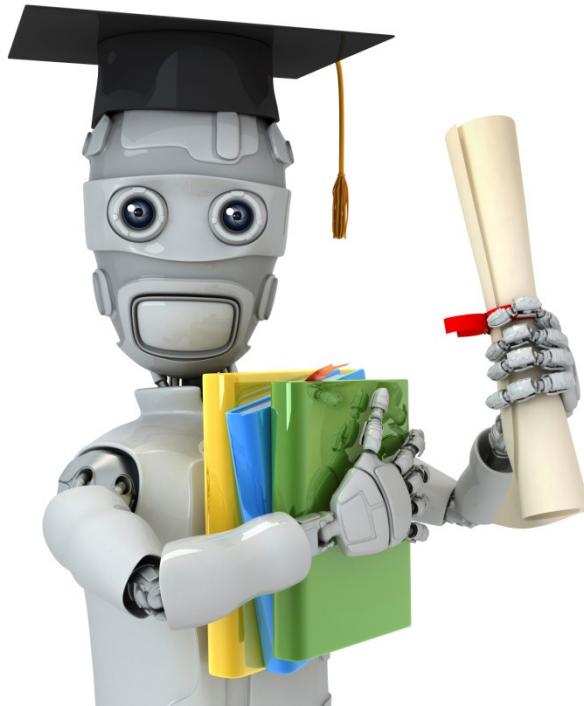
→ Problem 1: You have a large inventory of identical items. You want to predict how many of these items will sell over the next 3 months.

→ Problem 2: You'd like software to examine individual customer accounts, and for each account decide if it has been hacked/compromised.

→ 0 - not hacked
→ 1 - hacked

Should you treat these as classification or as regression problems?

- Treat both as classification problems.
- Treat problem 1 as a classification problem, problem 2 as a regression problem.
- Treat problem 1 as a regression problem, problem 2 as a classification problem.
- Treat both as regression problems.

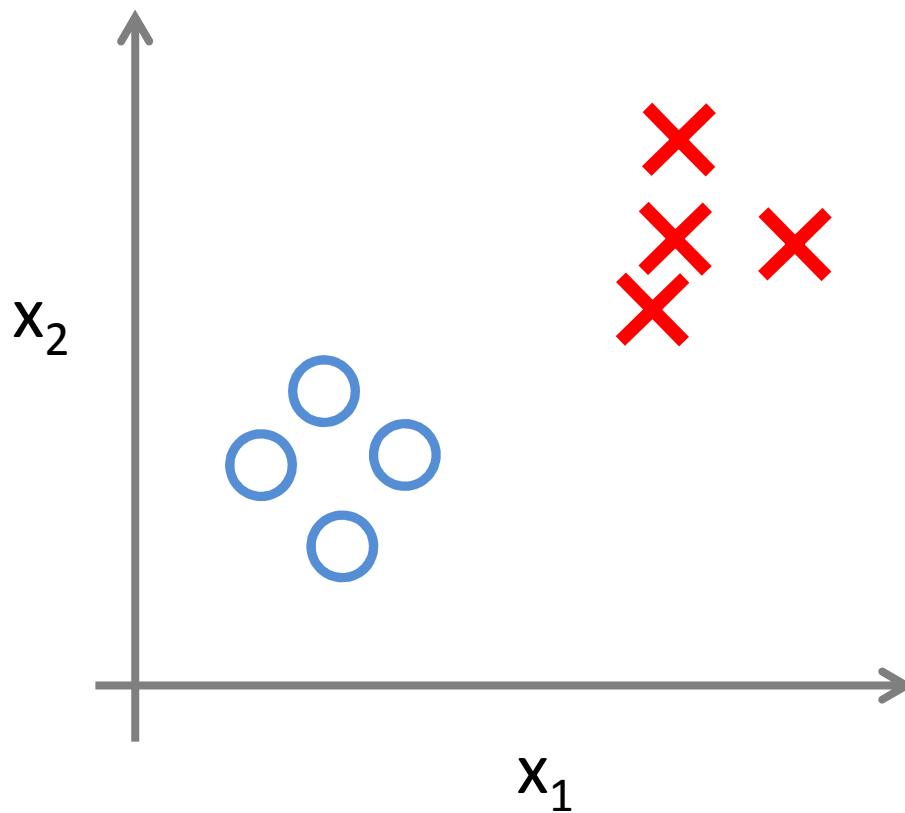


Machine Learning

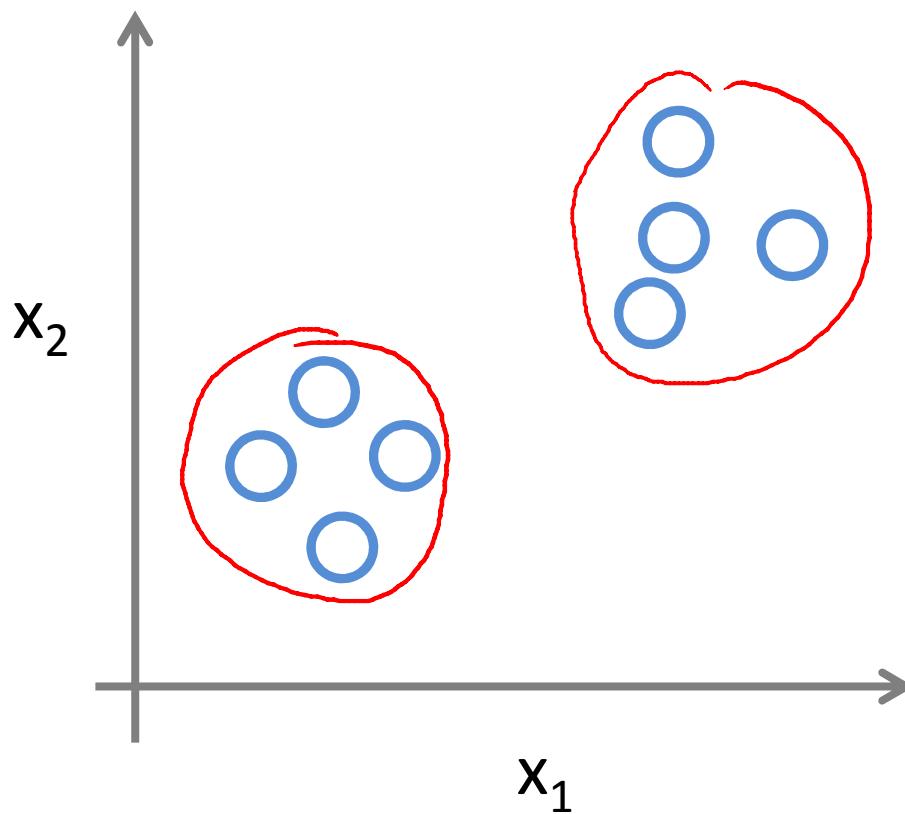
Introduction

Unsupervised Learning

Supervised Learning



Unsupervised Learning



Google News X

news.google.com X

Web Images Videos Maps News Shopping Gmail more ▾ andrewyantakng@gmail.com | Web History | Settings ▾ | Sign out

Google news

Search News Search the Web Advanced news search U.S. edition ▾ Add a section »

Top Stories

- Deepwater Horizon
- Fed meeting
- Foreign exchange market
- Lindsay Lohan
- IBM
- Tom Brady
- Toronto International Film Festival
- Paris Hilton
- Iran
- Hurricane Igor

Starred ★

- San Francisco Bay Area
- World
- U.S.
- Business
- Sci/Tech
- More Top Stories
- Spotlight
- Health
- Sports
- Entertainment

All news Headlines Images

Top Stories

[Christine O'Donnell »](#)
White House official denies Tea Party-focused ad campaign
CNN International - Ed Henry - 1 hour ago
Democratic sources say the White House is not considering an ad campaign tying Republicans to the Tea Party. Washington (CNN) -- A top White House official sharply denied a report that claims President Obama's political advisers are weighing a national ...
[Tea Party is misplacing the blame, former President Bill Clinton claims](#)
New York Daily News
[GOP tea party backer defends Christine O'Donnell](#) The Associated Press
Atlanta Journal Constitution - Politics Daily - MyFox Washington DC - Salon all 726 news articles »

US Stocks Climb After Recession Called Over, Homebuilders Gain
MarketWatch - Kristina Peterson - 16 minutes ago
NEW YORK (MarketWatch) -- US stocks climbed Monday, gaining speed after a key nonprofit organization officially called the recession over, giving investors a boost of confidence in the gradual economic recovery.
[Longest recession since 1930s ended in June 2009, group says](#)
Los Angeles Times
[Downturn Was Longest in Decades, Panel Confirms](#) New York Times
Wall Street Journal - AFP - CNN - USA Today all 276 news articles »

[Deepwater Horizon »](#)
BP Oil Well, Site of National Catastrophe, Dies at One
Vanity Fair - Juli Weiner - 22 minutes ago
The BP oil well, site of the Deepwater Horizon explosion that led to the worst oil spill in US history, died today at one year old.
Video: Blown-out BP Well Finally Killed in Gulf YouTube The Associated Press
Weiss Doubts BP Would End Operations in Gulf of Mexico: Video Bloomberg
CNN International - Wall Street Journal (blog) - The Guardian - New York Times all 2,292 news articles »

Recent

[Recession officially ended in June 2009](#)
CNNMoney - Chris Isidore - 39 minutes ago

[Hurricane Igor lashes Bermuda](#)
USA Today - Gerry Broome - 5 minutes ago

['Explain what you want from us,' reads front-page editorial](#)
msnbc.com - Olivia Torres - 10 minutes ago

Crisis response: Pakistan floods

San Francisco Bay Area - Edit

[Clorox »](#)
[Bay Biz Buzz: Clorox close to selling STP, Armor All](#)
San Jose Mercury News - 48 minutes ago - all 24 articles »

[Google's official beekeeper keeps the company buzzing with excitement](#)
San Jose Mercury News - Bruce Newman - 1 hour ago

[Jon Sylvia »](#)
[Martinez man still unconscious as police investigate weekend shooting](#)
San Jose Mercury News - Robert Salonga - 48 minutes ago - all 6 articles »

Spotlight

[Sarkozy rages at EU 'humiliation'](#)
Financial Times - Peggy Hollinger - Sep 16, 2010

news.google.com

Google news

Search News Search the Web Advanced news search

Top Stories

- Deepwater Horizon
- Fed meeting
- Foreign exchange market
- Harvey Weinstein
- Hillary Clinton
- IBM
- Tom Brady
- Toronto International Film Festival
- Paris Hilton
- Iran
- Hurricane Igor
- San Francisco Bay Area
- World
- Business
- Sci/Tech
- More Top Stories
- Spotlight
- Health
- Sports
- Entertainment

All news Headlines Images

Recent U.S. edition ▾ Add a section ▾

Christine O'Donnell's Home Office official denies Tea Party-focused ad campaign - CNN International - Ed Henry - 1 hour ago

Democratic sources say the White House is not considering an oil campaign type of proposal to combat climate change (CNHI) - A top White House official denied a report that claims President Obama's political advisers are weighing a national ... - TechCrunch - 1 hour ago

GOP tea party backer defends Christine O'Donnell - The Associated Press - AP Wire - 1 hour ago

New York Daily News - The blame, former Senator Bill Clinton claims

US Stocks Climb After Recession Called Over, Homebuilders Gain - MarketWatch - Kristina Peterson - 16 minutes ago

RECESSION - MyFox Philadelphia - 16 minutes ago

San Francisco Bay Area - Edit

Deepwater Horizon's BP Oil Well, Site of National Catastrophe, Dies at One - Vanity Fair - Juli Weiner - 22 minutes ago

BP's Deepwater Horizon oil well explosion that led to the world's largest oil spill in history, died at one year old - Video: Blown-out BP Well Finally Killed in Gulf - The Associated Press - AP Wire - 49 minutes ago

Who Dents BP? Would End Operations in Gulf of Mexico - Video: Bloomberg - CNN Wire - 49 minutes ago

Allen: Well is dead, but much Gulf Coast work remains

What next for Gulf oil spill?

STORY HIGHLIGHTS (CNN) -- The ruptured Macondo well, a mile under the Gulf of Mexico off the Louisiana coast, has been pronounced dead.

BP Kills Macondo, But Its Legacy Lives On - THE SOURCE

THE WALL STREET JOURNAL

Financial Services Transport Leisure Insurance Oil & Gas Sport Caught on the Web Betting Technology

SEARCH

BP Kills Macondo, But Its Legacy Lives On

Article Comments (2)

Email Print Permalink Like 2 Text +

By James Herron

BP confirmed late Sunday that the Macondo well that leaked almost five million barrels of oil into the Gulf of Mexico has been permanently sealed, but the well will continue to affect BP and the wider oil industry for many years.

The most immediate worry for BP and its shareholders is how the authorities will apportion blame for the spill. BP's own investigation spread responsibility across

Fire boat response crews battled the blazing remnants of the offshore oil platform on April 21, 2010. (Associated Press)

Most Recent

- Who Needs Plaza II Anyway
- Will Banks Be Forced to Split Retail And Banking Arms?
- Timing of Ratings Award Intriguing
- BP Kills Macondo, But Its Legacy Lives On
- We Already Need a Sental to Basal III!

Allen: Well is dead, but much Gulf Coast work remains

By the CNN Wire Staff September 20, 2010 — Updated 1317 GMT (2117 HKT)

Click to play

What next for Gulf oil spill?

STORY HIGHLIGHTS (CNN) -- The ruptured Macondo well, a mile under the Gulf of Mexico off the Louisiana coast, has been pronounced dead.

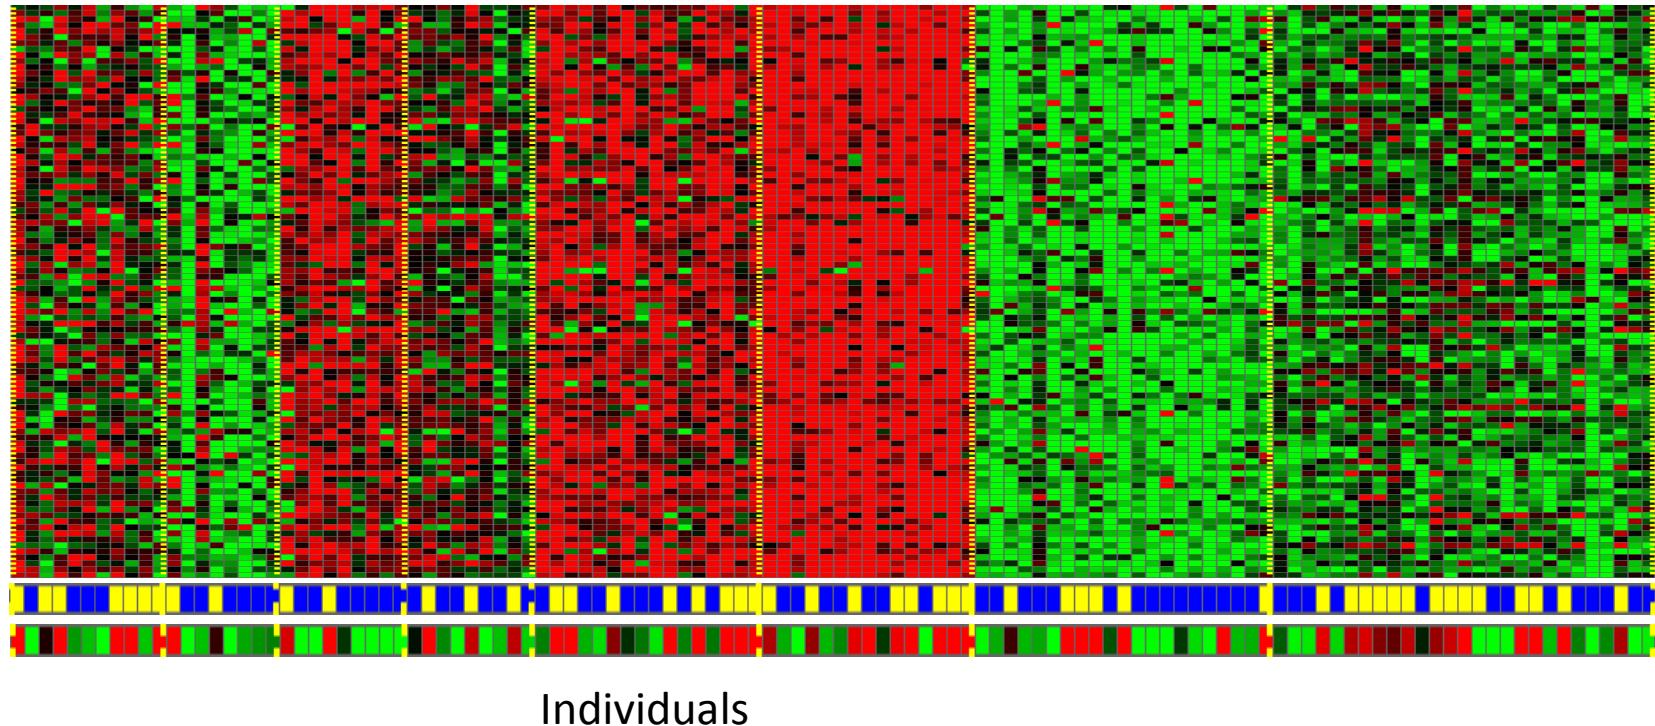
BP oil spill cost hits nearly \$10bn

BP has set up a \$20bn compensation fund after the Deepwater Horizon disaster, which has so far paid out 19,000 claims totalling more than \$240m

Julia Kollom guardian.co.uk, Monday 20 September 2010 08.33 BST Article history

BP's costs for the Deepwater Horizon disaster have hit \$10bn. Photograph: Ho/Reuters

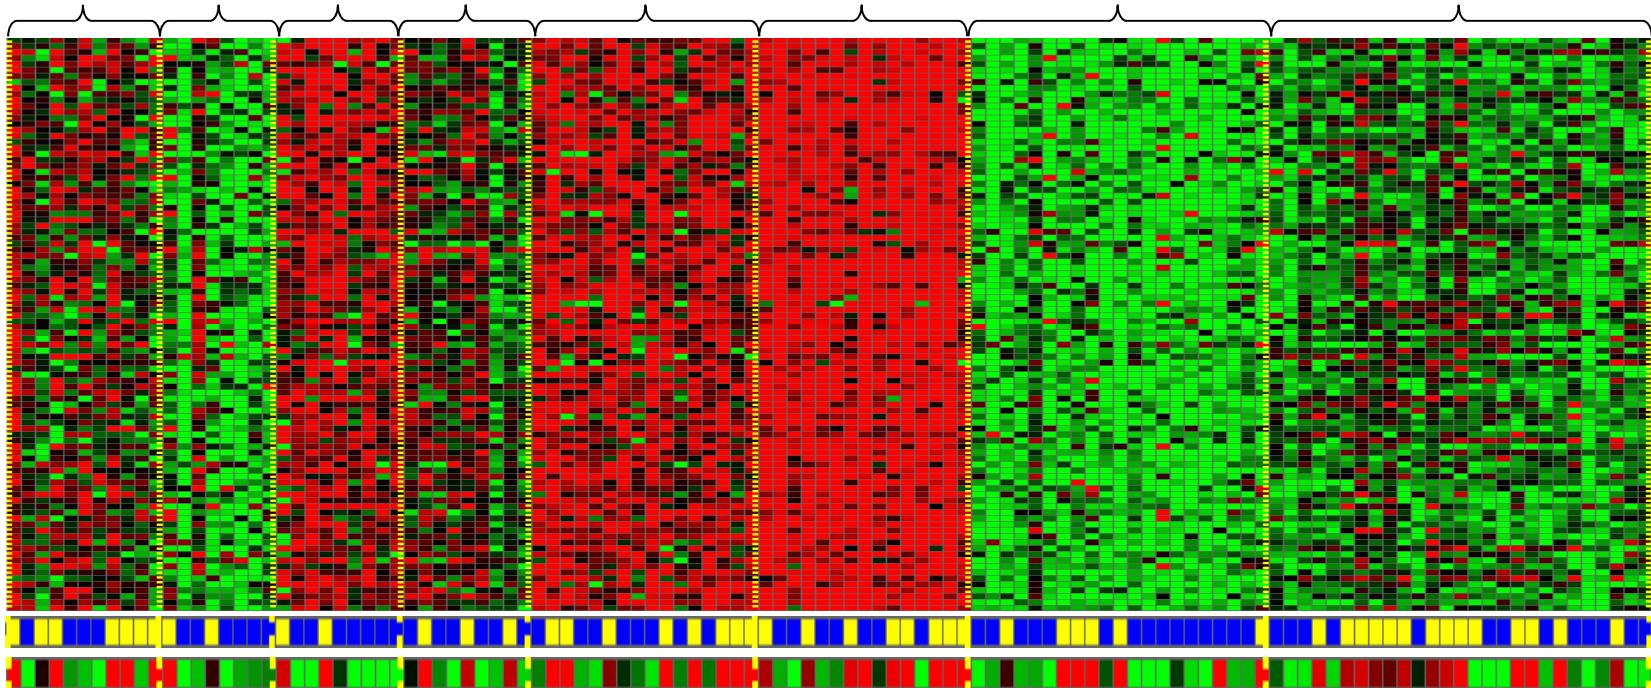
Andrew Ng



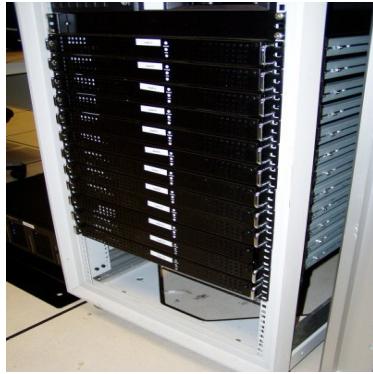
[Source: Daphne Koller]

Andrew Ng


Genes



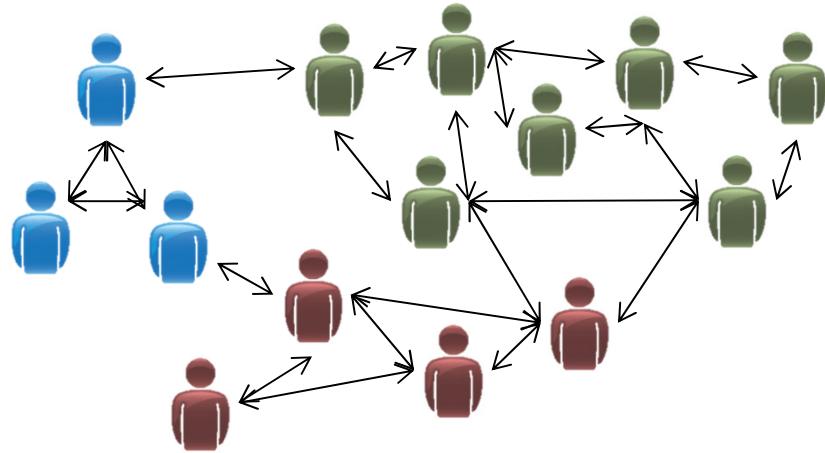
Individuals



Organize computing clusters



Market segmentation



Social network analysis

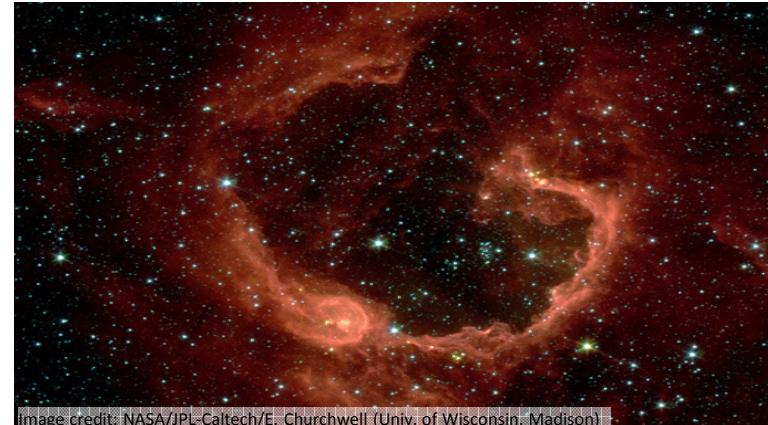
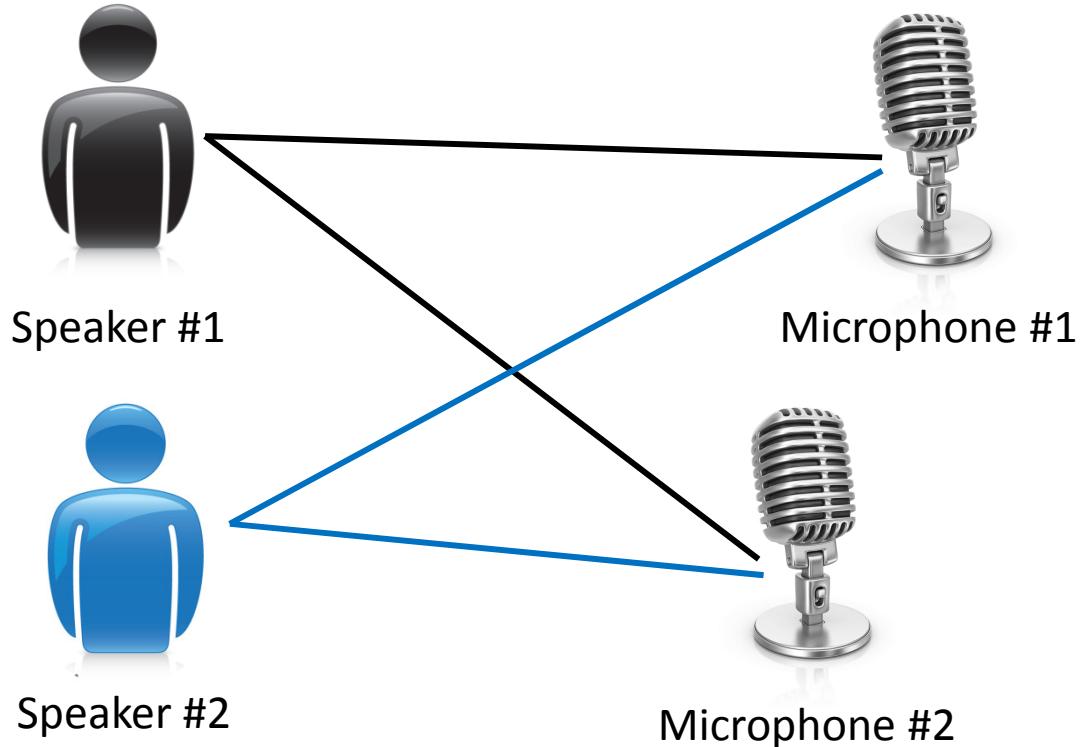


Image credit: NASA/JPL-Caltech/E. Churchwell (Univ. of Wisconsin, Madison)

Astronomical data analysis

Andrew Ng

Cocktail party problem



Microphone #1: 

Output #1: 

Microphone #2: 

Output #2: 

Microphone #1: 

Output #1: 

Microphone #2: 

Output #2: 

Cocktail party problem algorithm

```
[W,s,v] = svd((repmat(sum(x.*x,1),size(x,1),1).*x)*x');
```

Of the following examples, which would you address using an unsupervised learning algorithm? (Check all that apply.)

- Given email labeled as spam/not spam, learn a spam filter.
spam/not spam
- Given a set of news articles found on the web, group them into set of articles about the same story.
- Given a database of customer data, automatically discover market segments and group customers into different market segments.
- Given a dataset of patients diagnosed as either having diabetes or not, learn to classify new patients as having diabetes or not.

